NORMBANK: A Knowledge Bank of Situational Social Norms

Caleb Ziems 🌲😎 **Jane Dwivedi-Yu** ∞ **Yi-Chia Wang** ∞

Alon Y. Halevy ∞ **Diyi Yang** 🌲

🌲Stanford University, ∞Meta AI

{cziems, diyiy}@stanford.edu, {janeyu, yichiaw, ayh}@fb.com

Abstract

We present NORMBANK, a knowledge bank of 155k situational norms. This resource is designed to ground flexible normative reasoning for interactive, assistive, and collaborative AI systems. Unlike prior commonsense resources, NORMBANK grounds each inference within a multivalent sociocultural frame, which includes the setting (e.g., *restaurant*), the agents' contingent roles (*waiter*, *customer*), their attributes (*age*, *gender*), and other physical, social, and cultural constraints (e.g., the *temperature* or the *country of operation*). In total, NORMBANK contains 63k unique constraints from a taxonomy that we introduce and iteratively refine here. Constraints then apply in different combinations to frame social norms. Under these manipulations, norms are *non-monotonic* — one can cancel an inference by updating its frame even slightly. Still, we find evidence that neural models can help reliably extend the scope and coverage of NORMBANK. We further demonstrate the utility of this resource with a series of transfer experiments. For data and code, see https://github.com/SALT-NLP/normbank

1 Introduction

As AI systems continually evolve for human assistance and collaboration, they will increasingly operate within cultural and social spaces, and require increasingly robust and flexible knowledge of social norms (Carlucci et al., 2015). From dialogue systems (Molnár and Szüts, 2018; Vaidyam et al., 2019; Bavaresco et al., 2020; Grossman et al., 2019) to socially interactive robots (Fong et al., 2003; Deng et al., 2019) and augmented or mixed reality technologies (Anderson and Rainie, 2022), each could benefit from understanding how humans effectively communicate, make decisions, engage with requests, and broadly interact with others (Sunstein, 1996; Sherif and Sherif, 1953).
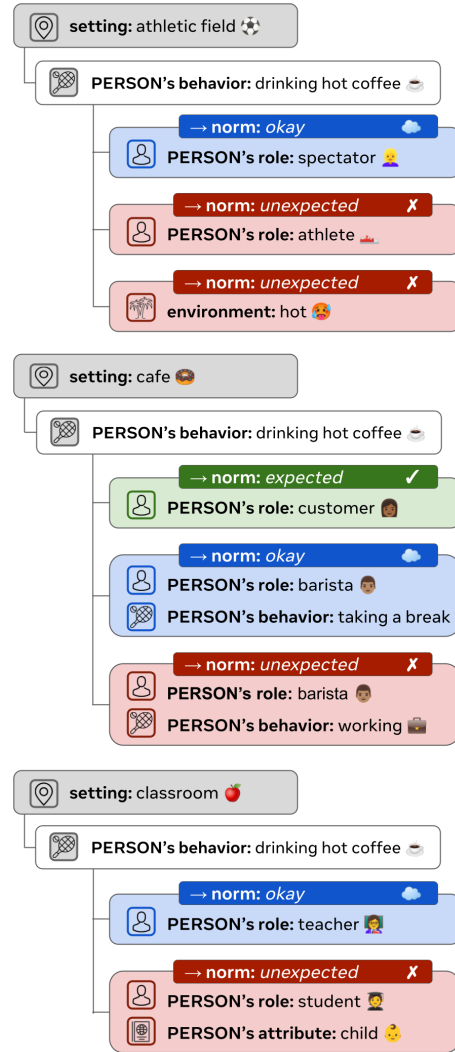
😎 Work done at Meta AI Research

Figure 1: **What is special about NORMBANK?** Norms are grounded by *situational constraints*—environmental and personal attributes, as well as roles and other behaviors. In this example, drinking coffee is an encouraged activity in its prototypical context, for a *customer* in a *cafe*, but it is counternormative for a working barista to do so in the same cafe, or for a child-age student do so in a classroom. These represent only some of the non-monotonic normative inferences that are represented in NORMBANK.

Natural language is flexible and highly expressive; thus it is a promising medium for encoding knowledge of social norms (Sap et al., 2019a). The goal of this work is to construct NORMBANK, a natural language bank of social norms that will allow AI systems to reason about social situations under complex constraints. NORMBANK encodes 155k norms via scalable human annotation, bootstrapped with implicit knowledge from large pre-trained language models (LLMs).

NORMBANK factors in two important considerations that have been previously overlooked. First: norms are not rigid truths; they are flexibly assumed standards that may be updated with new information from the social context (Blass and Horswill, 2015). Second: this social context is not a flat list of facts but a matrix of hierarchical dependencies (Hovy and Yang, 2021). These two considerations have important design implications for norm representations and reasoning in AI, inspiring two objectives for this work.

Objective 1 is to support *non-monotonic* reasoning (Reiter, 1981) over *defeasible* (Pollock, 1987) norms. This means inferences that hold under most cases can be updated or even retracted based on new information. For example, *dancing* is a positive behavior that is generally permitted in many casual settings and in many cultures. We can still strengthen or cancel this inference. On the one hand, dancing is expected from a professional dancer. But in an Islamic cultural context, individuals are forbidden from publicly dancing with members of the opposite sex. In a hospital setting, a young child is allowed to dance in the waiting room, but this behavior would not be expected from an adult visiting a dying relative. For more examples, see Figure 1. This kind of reasoning will not always follow straightforward compositional logic (Klimczyk, 2021), and we expect it to be a challenge for AI systems.[1] NORMBANK is the first data resource to support non-monotonic normative reasoning by encoding *contrasting* situations under which the *same behavior* could alternatively be expected or considered taboo (see §4).

There is a combinatorially explosive space of situational contexts, each with non-compositional and thus unpredictable norms. Enumerating the set of all possible constraints is intractable. To effi-

ciently learn norms in this space, models can rely on the regularizing effects of hierarchical organization and social theory. NORMBANK introduces hierarchical organization (**Objective 2**) by means of a rich taxonomy over the relevant contextual signals that inform behaviors.

Our new SCENE taxonomy is the first to use Goffman's (1959) dramaturgical theory of social life. We operationalize the theory with *settings* that have additional *environmental* constraints. In each setting, there are agents with different *roles* and *attributes*, who then perform *behaviors*. Norms apply to behaviors in certain situations. For example, in Figure 1, norms around *drinking hot coffee* differ for agents with different roles (e.g., *barista, customer*) and attributes (e.g., *adult, child*), in different settings (e.g., *cafe, classroom*).

Having addressed the objectives above, we train neural models to expand NORMBANK with automatic knowledge completion. Experiments show promising results: these models can extrapolate social commonsense to new behaviors in new situations, leveraging similarities in analogous roles across different situations. Finally, we demonstrate how to transfer knowledge via sequential finetuning from NORMBANK to social reasoning tasks. Together, knowledge completion and transfer learning suggest that our dataset will serve as a useful resource for advancing neural models toward situationally-grounded social reasoning.

2 Related Work

Commonsense knowledge bases (CSKBs) are sets of structured knowledge about everyday life. They capture broad taxonomic relationships (Liu and Singh, 2004; Speer et al., 2017; Elsahar et al., 2018), logical relations (Lenat, 1995; Zhang et al., 2018), and universal laws of causality and physical mechanics (Talmor et al., 2019; Bisk et al., 2020). More recent datasets encode *social mechanics*, like broad human values (Ziems et al., 2022; Hendrycks et al., 2021), norms (Forbes et al., 2020; Fung et al., 2022), and typical rules of social behavior and motivation (Sap et al., 2019a; Huang et al., 2019). NormBank specifically centers social norms around dramaturgical settings (i.e., *places of worship, commerce, and recreation*). Just as ATOMIC (Sap et al., 2019a) seeded Social IQa (Sap et al., 2019b) and δ-NLI (Rudinger et al., 2020), we anticipate that NormBank can be converted into benchmarking tasks, plus injected into language models

---

[1]Typically, it's okay to drink soda while actively working and it's okay for a waiter to drink soda; yet the intersection of these conditions is not typical. It is **NOT** okay for a waiter to drink soda while actively working.

for downstream applications (Chang et al., 2020; Mitra et al., 2019; Ji et al., 2020a,b).

Norm discovery is an emerging method inspired by automatic knowledge base construction (Mitchell et al., 2015; Weston et al., 2013; Craven et al., 2000) and extracting social knowledge from LLMs via prompting (Trinh and Le, 2018; Petroni et al., 2019; Wang et al., 2019; Sakaguchi et al., 2020). In concurrent work, Fung et al. (2022) propose NORMSAGE, which automatically discovers mandated or conventional behaviors from dialogues. Their prompts resemble our bootstrapping efforts in §3, with the added step of automatic self-verification. NORMBANK differs from NORM-SAGE in that we rely on human annotation to collect more creatively non-prototypical situations to challenge and expand normative reasoning models.

Normative reasoning systems like Delphi (Jiang et al., 2021), and UNICORN (Lourie et al., 2021a) are pre-trained on existing social knowledge bases (Forbes et al., 2020; Emelin et al., 2021; Hendrycks et al., 2021; Lourie et al., 2021b; Sap et al., 2020), which contain more conventional social behaviors from narrative contexts. Until Pyatkin et al. (2022), in work concurrent to our own, descriptive social reasoning systems have been framed as universal oracles with forced-choice judgments about human behaviors (Talat et al., 2022). These models lack the capacity for defeasible reasoning (Madaan et al., 2021; Rudinger et al., 2020). Oracles instead tend to assume the most prototypical contexts (Boratko et al., 2020). Many of these predictions will appear reasonable if we pragmatically infer a conventional narrative, but for systems to achieve robust social intelligence, they must account for the long tail of the distribution. We can easily find unconventional contexts in which the correct inference contained in NORMBANK is misunderstood by current models.[2]

3 SCENE: A Dramaturgical Framework

> *The self... is a dramatic effect arising diffusely from a SCENE.* 🎬
> — **Erving Goffman** (1959)

To help models efficiently learn non-monotonic normative reasoning over a seemingly unbounded

set of possible contexts, and to test this understanding in LLMs, we will need to establish a more tractable set of elements to represent this social matrix. For this purpose, we construct a hierarchical taxonomy of constraints, which we call *Situational Constraints for social Expectations, Norms, and Etiquette* (SCENE for short). SCENE follows Goffman's (1959) dramaturgical model of social life. According to this model, people are like actors trying to maintain a social performance in front of an audience. Each actor performs a particular *role* as if in a scene from a movie. The scene is grounded in a particular *setting*, which includes aspects of the *environment* that inform the performance. Each scene also has a script (Schank and Abelson, 1977), which tells the actor what kinds of *behaviors* will be perceived as in-character or out-of-character. Additionally, the actor will embody socially meaningful *attributes* like age, gender, status, etc. These attributes may be relevant to the scene and the actors place in it. In Figure 2, the example setting is a *restaurant* where the environment is *uncrowded* and the hour is *night*. There are two primary roles of *customer* and *server*, and for norm formation, some relevant attributes include their respective *genders*, *sexualities*, and *ages*, which parameterize the behaviors that are appropriate for this dinner, such as *dating* and *drinking alcohol*.

Settings ◎ (e.g., banks, classrooms, homes, hospitals) are the loci of scripted social interactions (Schank and Abelson, 1977), and they frame all subsequent elements of NORMBANK, so we begin here with 129 distinct settings like *amusement park, bus,* and *elevator*. Settings derive from two popular knowledge resources. First, there are 80 settings from from ConceptNet (Speer et al., 2017), a broad knowledge base of the words and phrases that people commonly use.[3] There are another 255 settings from the "movie scene" label in the MovieGraphs (Vicol et al., 2018) resource—a collection of social situations that were depicted in movie clips.

The Environment 🎋 contains signals that can trigger associative priming of social norms (e.g., the noise level of a study space; Aarts and Dijksterhuis). This portion of the taxonomy is designed to be broad and general-purpose, with a set of attributes that can refine any setting. Our taxonomy is based on a broad review of the literature on norm

---

[2]For example, Delphi believes *yelling* and *clenching your fists*, *breathing heavily*, or *asking someone personal questions about their sex life* are all conventionally inappropriate. NORMBANK gives acceptable contexts for each: *guests riding a roller coaster, athletes running track, and doctors performing routine checkups,* respectively.

[3]Settings were defined by head-entities with an IsA relationship to some tail in {place, location, area}. Manual inspection proved the usefulness of this heuristic.
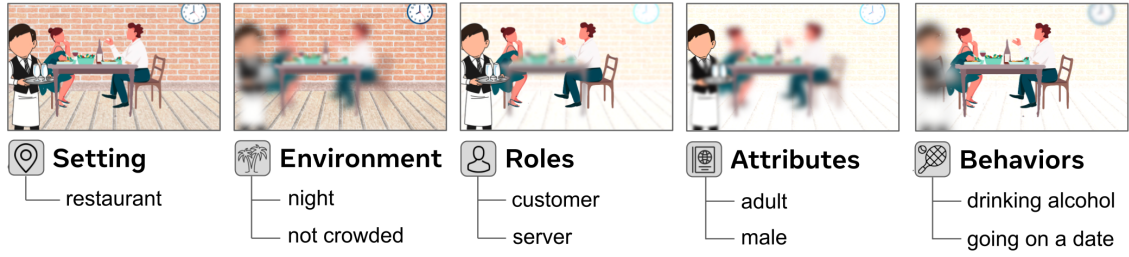
Figure 2: **An example of the SCENE Dramaturgical Framework used to constrain NORMBANK**. The *restaurant* setting is specified by the **attendance** (*not crowded*) and **time of day** (*night*) in the environment. The two agent roles, *customer and server*; the latter is specified by the **age bracket** (*adult*) and **gender** (*male*) attributes. The former are engaged in the behaviors *drinking alcohol* and *going on a date*. *Note:* Graphics are for illustration. NORMBANK is a text dataset and does not contain any images.

formation and its relevant factors (van Rijswijk and Haans, 2018; Janicik and Bartel, 2003; Boyce et al., 2000; Russell and Ward, 1982; Durkheim, 1915). Importantly, the taxonomy is further refined through crowdsourced feedback (§4). Ultimately, our taxonomy grows to contain 404 environmental constraints. An extensive overview of the environmental constraints is given in Appendix B.1, but we summarize them here.

In the environment, there are important taxonomic subclasses of factors that inform norms. One subclass is **time** constraints, like seasonality (Janicik and Bartel, 2003), holidays and special customary observances (Durkheim, 1915), and another is the **country of operation**, which serves as a proxy for regional cultural differences (Meyer, 2014). We also include factors from **environmental psychology** (Bell et al., 2001) that involve the agent's comfort and ease in the environment (e.g., noise level, privacy, and cleanliness). Additionally, **physical conditions** include factors like weather, which impact visibility, coordination, safety, and comfort (Boyce et al., 2000; van Rijswijk and Haans, 2018; Cunningham, 1979). In addition to the imposed taxonomy, annotator feedback (§4) lead us to add a subclass called **restrictions** that formally limit attendance, participation, and behavior, due to notions of formality, religiosity, or exclusiveness.

Roles 👤 may be ubiquitous, but it is challenging to collect reliable, setting-specific roles with high coverage. Our solution is to use the powerful associative knowledge of LLMs to automatically enumerate roles for each setting via prompting, inspired by Trinh and Le (2018), Petroni et al. (2019), and others (Wang et al., 2019; Sakaguchi et al., 2020). Specifically, we prompt GPT-3 (Brown et al.,

2020) `text-davinci-002` in a zero-shot manner with the phrase "`Some roles <preposition> <determiner> <setting>:`" where the preposition and determiner are manually configured to match the setting (for example, "`some roles at a casino:`" or "`some roles on the beach:`"). On average, we generate 5.5 roles per setting, with a total of 928 unique roles.

Attributes 🌐 are properties of individual agents that determine their social norms. Here again, the goal is to derive a general purpose taxonomy from the literature. Some attributes are basic demographic categories like the person's **age bracket**, **gender**, **race**, **religion**, and **sexuality** (Thompson Jr and Pleck, 1986; Dempsey and De Vaus, 2004; Helgeson, 2016). Related demographic categories include **education** level, **employment**, and **marital status**. Since food is a focal point for culture and morality, we include **diet**. We also include material constraints like **medical condition** and **social class**. Finally, we increase the coverage of this set by including generic descriptors of two types: **condition or state** adjectives, which describe a temporary condition (e.g., *dizzy*), and **characteristic** adjectives that describe more permanent attributes (e.g., *blonde*). In total, our taxonomy defines 578 attribute constraints.

Behaviors 🎲 are the primary target of analysis for social norms. As with roles, we co-opt GPT-3 to enumerate behaviors for each setting and role, but the approach here is augmented in two ways. First, we include a norm *expectation* in the prompt. By querying for *unexpected* behaviors, we can begin to shift the distribution of behaviors away from the prototypical. Second, we increase the diversity of generations by conditioning on

the agent's attribute. This further reduces the number of conventional behaviors in our set. The prompt is `"Some things you would (never) do <preposition> <determiner> <setting> (if you were <attribute>):"` where elements in parentheses are optional elements. In this way, we generated an average of 776.5 behaviors per setting, which was filtered down to 112.6 behaviors per setting, via programmatic methods described in Appendix B.2.

4   Building NORMBANK

Section 3 gave us a high-recall set of constraint variables for explaining situational social norms. Our end goal is to build a resource that contains reliable norms to ground, train, and test automatic normative reasoning systems. We want these norms to describe challenging, non-prototypical examples, and to depend on subtly contrasting situations that, when shifted, change the norm label non-monotonically. This motivates us to use human annotation over the rich SCENE taxonomy.

Our process is essentially the reverse of the current paradigm established by prior work, which starts with a basic narrative context and subsequently extracts (Fung et al., 2022) or annotates (Forbes et al., 2020) the expected behaviors. Instead, we start with behaviors and ask annotators to provide us with different dramaturgical contexts (SCENE constraints) under which that behavior could be variously seen as *expected*, *okay*, or *unexpected*. Thus we obtain richer and less prototypical instances—examples not mentioned in standard dialogue, which will significantly challenge models. The approach is inspired by contrast sets (Gardner et al., 2020) and counterfactual augmentation (Kaushik et al., 2020) as means of reducing spurious correlations in model inferences.

4.1   Annotation Task

For the annotation task, we recruit experienced English-speaking Mechanical Turk annotators who have $\geq$98% acceptance with $\geq$100 HITs and are located in the United States. The task requires human creativity over a large combinatorial space. For a given setting s and a behavior b, an annotator will tell us distinct situational contexts under which b is alternatively *expected* (required by duty or anticipated with high probability), *okay* (permitted or anticipated with moderate probability), or *unexpected* (forbidden, stigmatized, taboo, or otherwise

anticipated with very low probability).

These *expected, okay,* and *unexpected* categories are called "norm labels." The language of expectation is useful for describing behavioral regularities—the focus of this work—rather than enumerating top-down or bottom-up judgements of *ethical* or *moral* behavior, as in prior datasets (Ziems et al., 2022; Emelin et al., 2021; Lourie et al., 2021b). Importantly, we do not impose any ethical philosophy or framework as in Hendrycks et al. (2021), but instead, encourage annotators to find norms that merely describe observable social life (Cialdini et al., 1991).

The annotator fully specifies the appropriate situational context by means of disjunctions and conjunctions of constraints. For example, "spit at a dentist's office" can be unexpected when `(PERSON's role is 'dentist')` or when `((PERSON's role is 'patient') AND (PERSON's behavior is 'checking in'))`. Annotators select SCENE constraints using drop down menus that follow the hierarchy of §3 (for details on the HIT interface, see Appendix C.2). They are also free to insert their own custom constraints into the hierarchy. In this way, we iteratively expand the taxonomy.

4.2   Dataset Quality

Quality Control. Manual inspection of over 2.5k data points reveals that the open-ended and creative aspects of the task are natural incentives for high-quality work (Chandler et al., 2015; Sheehan, 2018). To further ensure the quality of NORM-BANK, we trained annotators with careful instruction, a qualification test, a staging round, personalized feedback, programmatic filtering, and finally, a series of random audits (Litman et al., 2015; Sheehan, 2018). The instructions included at least 3 fully-worked examples for each norm label, plus suggestions and explanations for a total of 24 constraints. We administered a six-question qualifier, which tested workers' knowledge of the taxonomy, definitions, free text response, and how to properly indicate constraint conjunctions and disjunctions through the task interface. If the worker passed at least five questions correctly on the first try, she would gain access to the staging round – a small-scale version of the task in which each submission would receive detailed and personalized feedback.

We invested a significant amount of time to feedback, offering 75 to 200 words of review for each

	Constraints	Norms	Situations	Constr. / Norm	Taxonomic Constr.	Pre-populated Constr.
train	328,045	124,920	57,417	2.63	93.5%	69.3%
dev	37,761	15,008	8,573	2.52	92.8%	66.8%
test	42,601	15,495	8,674	2.75	94.6%	70.9%
NORMBANK	408,407	155,423	70,215	2.63	93.6%	69.2%

Table 1: **Summary statistics** show the immense scale of NORMBANK (§4) and the broad coverage of our SCENE framework (§3). There are 155k total annotated norms, comprised of 70k unique situations, and each situation is drawn from a conjunction of some subset of the 408k annotated constraints. Of these annotated constraints, 94% of them use the structure of our SCENE taxonomy, and 69% use a pre-populated constraint value from one of our taxonomic dropdown menus.

of 2,502 staging HITs. Once a worker submitted 3 high-quality HITs in the staging round, he or she could move to the full task. To identify poor work here, we programmatically flagged workers with extremely low variation in their annotations. Finally, we periodically performed a total of three random audits, sampling 250 annotations in each audit, to confirm the quality of the annotations. Workers were paid a base rate, plus an additional itemized bonus for every additional constraint they added, which incentivized workers to be more expressive and creative. Annotators received a median of $30 per hour for this task.

Quality Metrics. The above methods all proved remarkably successful in generating a creative and high-quality resource. Because our task is creative and subjective, data quality is not easily measured by inter-annotator agreement. We instead report human evaluations over the Gold NORMBANK data in the bottom row of Table 3 (alongside model generations from §5.1). Annotations are considered sensible (82.5%), relevant (82.17%), and normative (72.9%). Still, it is important to note that around half of gold annotations are marked by third-party evaluators as fully correct by majority vote.

With regard to the correctness metric, annotator disagreements can be traced to differences in the annotators' models of the world, which likely stem from their own personal differences, including age, profession, and worldview. For example, an annotator likely familiar with the Cambodian tradition of "Pithi Srang Preah" marked that "honoring your ancestors" is normal for Cambodians on Cambodian New Year, while an annotator unfamiliar with this practice marked it as unexpected. Furthermore, we administered political leaning and the moral foundations surveys to all annotators, which we release alongside NORMBANK to help explain how these
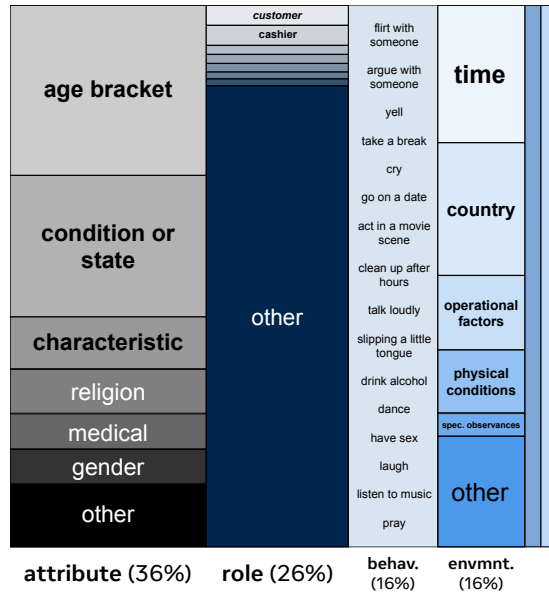

Figure 3: **Distribution of NORMBANK constraints** where the area of each cell is proportional to the distribution. Constraints are dominated most by the agent's attributes and roles, with a smaller and even split between behaviors and the environment.

personal differences informed the probabilities they assigned to events. This resource will be of interest to computer scientists and social scientists, since NORMBANK contains not only commonsense facts, but also culturally-conditioned distributions over behavior and expectations about behavior.

4.3 Dataset Summary

Summary Statistics. Table 1 gives the summary statistics for the annotated dataset. NORMBANK contains a total of 408k constraints, applied to 155k norms, for an average of 2.63 constraints per norm. The SCENE taxonomy broadly captures the kinds of constraints annotators were looking for 94% of

	Norm Classification Results			
Model	P	R	F	Acc
ALBERT	68.0	66.6	67.1	71.0
BERT	72.1	70.7	71.3	74.6
RoBERTa	**73.3**	**71.4**	**72.1**	**75.4**

Table 2: **Classification Results** for {*expected, okay,* or *unexpected*} show that standard transformer models can learn to make normative inferences with an accuracy that is adequate for expanding NORMBANK.

the time, and they were able to find their exact constraint value from a pre-populated list in 69% cases. For concurrent behavior and attribute constraints, annotators had to input their own values in 59% and 33% of cases respectively, followed by 27% and 9% of cases for the environment and roles. Overall, this indicates that our GPT-3 prompting method achieved high recall, especially for roles, and least so for behaviors, which is unsurprising, given the almost unbounded space of viable human behavior.

Figure 3 gives the distribution of constraints in NORMBANK. Constraints are dominated most by the agent's attributes and roles. Age, condition, and characteristic are the most popular attributes, while roles vary. There is an even split between behaviors and the environment. In the environment, there is a notable focus on *time* constraints, and slightly lesser but more even attention towards the remaining subcategories.

Links to Existing Knowledge. NORMBANK's SCENE taxonomy has close links to existing knowledge resources. ConceptNet directly seeded 80 settings in SCENE. Beyond this, we successfully link over 90% of taxonomic items from the setting, environment, roles, and attributes directly with concepts in ConceptNet. These taxonomic items cover 93.6% of all constraint categories and 70.0% of all constraint values. ConceptNet is further linked to WordNet, DBPedia, Umbel, Cyc, and Wiktionary, so by extension, NORMBANK can be coupled to these resources.

5 Experiments: How to use NORMBANK

NORMBANK is not designed for any particular narrow task; it is designed as a general-purpose knowledge resource that can ground social reasoning through downstream tasks (compare ATOMIC

(Sap et al., 2019a) and ConceptNet (Liu and Singh, 2004)). Towards this end, NORMBANK should contain richly organized knowledge that can be learned by neural models and applied for non-monotonic reasoning in new settings. In this way, it should be possible to automatically expand the NORMBANK resource. The knowledge contained here should also be applicable across a range of social reasoning tasks. Thus our experiments aim to demonstrate two things: (§5.1) that we can automatically expand NORMBANK using neural methods, and (§5.2) that NORMBANK is a useful resource with relevant knowledge for downstream applications. For all experiments in the following subsections, we use an 80%-10%-10% train-dev-test split in which <setting, behavior> tuples in one set are never seen in another.

5.1 Automatic Knowledge Completion

How can we expand NORMBANK? We considered two different methods of knowledge bank completion (Weston et al., 2013; Craven et al., 2000), which rely on different assumptions. Results from both methods indicate that NORMBANK is rich enough to support its own automatic expansion. Classification is the simpler case, where we assume a closed world (Bordes et al., 2013; Lin et al., 2015b,a; Socher et al., 2013), while generation assumes an open world (Shi and Weninger, 2018) with a modifiable set of constraints.

Classification. Here, our known constraints and behaviors (§3) will remain fixed (Shi and Weninger, 2018), but we can discover new relationships by classifying unseen behavior and constraint combinations as *expected, okay,* or *unexpected*. The advantage of this approach is that it is straightforward, and the disadvantage is that evaluating classifiers over the power set of the entire constraint space would be intractable; thus more efficient search methods will be needed.

For all norm classification tasks, we fine-tune three popular transformer models: BERT-base-uncased (Devlin et al., 2019), RoBERTa-base (Liu et al., 2019), and ALBERT-base-v2 (Lan et al., 2020), with hyperparameters in Appendix A. Results in Table 2 appear promising. Given the scope and scale of NORMBANK, models are capable of learning non-monotonic inferences, achieving F1 scores as high as 72.1% on the test set. This shows that classification is a reliable method for NORM-BANK knowledge expansion.

Model	Decoding	ROUGE-L	BLEU	Avg. # Constr.	Tax. Constr.	Pre-pop. Constr.	Sensible$_S$	Correct$_S$	Norm$_C$	Relevant$_C$
BART	greedy	75.7	23.1	1.62	93.8	40.7	**100.0**	46.0	**94.3**	94.3
	beam	**77.1**	23.5	1.70	98.8	48.2	98.0	36.0	92.5	94.9
	p=0.9	75.8	23.1	1.62	87.7	42.0	**100.0**	46.0	91.9	**97.0**
GPT-2	greedy	57.4	12.5	1.78	70.8	36.0	86.0	34.0	86.1	89.6
	beam	70.0	16.5	1.26	95.2	47.6	94.0	34.0	88.3	94.3
	p=0.9	60.2	13.2	1.46	90.4	43.8	84.0	34.0	93.0	94.0
T-5	greedy	45.0	10.3	2.70	67.2	28.4	60.0	22.0	78.8	78.9
	beam	68.3	13.7	1.02	100.0	64.7	86.0	40.0	89.5	96.3
	p=0.9	48.4	10.9	2.26	72.7	30.0	76.0	42.0	83.0	79.1
GPT-3 davinci-002		44.2	23.2	5.31	85.4	33.0	87.3	49.0	90.5	83.5
GPT-3 davinci-003		51.7	**28.6**	2.34	84.6	33.2	95.0	**61.1**	91.8	87.8
Gold NORMBANK		-	-	2.68	93.6	70.0	82.5	55.0	72.9	81.7

Table 3: **Constraint generation results.** (*Left*) Automatic evaluation suggests that BART has the advantage over other generative models. (*Middle*) Generated constraints fall into the SCENE taxonomy 67.2-100% [Tax. Constr.] and use a pre-populated constraint 30-64% of the time [Pre-pop. Constr.], depending on the decoding strategy. (*Right*) Human eval shows encouraging results: a NORMBANK-trained BART can generate sensible, correct, normative, and relevant constraints for use in automatically expanding NORMBANK. Here, the best fine-tuned model results are highlighted, while the best overall model results are **bolded**.

Generation. The model is trained with a forward language modeling objective over the string g:

$$g = \{[\texttt{SETTING}], s_1, s_2, ..., s_n,$$
$$[\texttt{BEHAVIOR}], b_1, b_2, ..., b_m,$$
$$[\texttt{NORM}], \texttt{label},$$
$$[\texttt{CONSTRAINTS}] c_1^1, c_2^1, ..., c_{\ell_1}^1$$
$$[\texttt{AND}] c_1^2, c_2^2, ..., c_{\ell_2}^2 ...$$
$$[\texttt{AND}] c_1^k, c_2^k, ..., c_{\ell_k}^k \texttt{<EOS>}\}$$

At inference time, the model generates the list of constraints $c^1, ..., c^k$ that will make the norm `label` true as it is conditioned on the setting s and behavior b. For this purpose, it is sufficient to use BART (Lewis et al., 2020), GPT-2 (Radford et al., 2019), and T5 (Raffel et al., 2020), three powerful language models used widely for generative inference. We also prompted GPT-3 davinci-002 and davinci-003 in a few-shot manner via the OpenAI API (see the prompts in Appendix A).

Evaluation comes from both automatic and human metrics. Humans evaluate 300 <setting, behavior> data points for each of the 11 Model × Decoding combinations, plus gold standard examples from NORMBANK. For constraints, they provide us the **% Norm$_C$** (*proportion that helps represent a human rule or expectation for behavior*) and **% Relevant$_C$** (*proportion that relates to the norm without redundancy or tautology*). They also give us situation[4] evals: are they **% Correct$_S$** (*they produce an accurate norm label*) and mutually **% Sensible$_S$** (*all constraints can be true at the same time*).

---
[4]Situations are defined as the intersection of constraints

Table 3 gives the generation results for the three fine-tuned models: BART, GPT-2, and T-5. According to human judgment, all models produce text that successfully constrains human expectations for behavior (*Norm$_C$* ~90%). BART + nucleus sampling ($p = 0.9$) gives the most *Sensible$_S$* (100%) and *Correct$_S$* situations (46%) with the most *Relevant$_C$* constraints (97%). This is clearly a challenging task: situations are deemed correct only 46% of the time. Yet they closely approach the scores of human gold-standard data (55%). Notably, generated constraints are highly relevant to the norm label and entirely mutually-sensible. Given the challenging nature of the task, the results are quite encouraging, suggesting that NORMBANK can facilitate its own expansion via natural language generation.

Prompting Results in Table 3 show that few-shot GPT-3 models fail to match our best performing BART model's ability to generate *Sensible* situations (95 vs 100%) with high *%Norm* (91.8 vs 94.3%) and *%Relevant* constraints (87.8 vs 97%). Still, annotators are more likely to find GPT-3 output to be *Correct* overall (61.1 vs 46%). Automatic metrics show that GPT-3 achieves higher precision (28.6 vs 23.1 BLEU) at the expense of recall (51.7 vs 77.1 ROUGE-L), suggesting that GPT-3's generations, while often correct, may be more prototypical. Qualitative analysis confirms this.

Sometimes the conventional answer leads GPT-3 astray, as when it uses a series of faulty lexical associations to explain that 'drinking milk' is *unexpected* on an 'athletic field' for individuals who are not the coach and for those whose behavior is not

'hydrate' while the temperature is 'warm.' BART, however, correctly discerns that it's unexpected for athletes whose behavior is 'playing sports.' In general, GPT-3 appears more likely to underspecify the situation. For example, GPT-3 responds that it's *expected* for a homeowner to 'leave the gate open' in the 'backyard,' and BART agrees, but BART further specifies that the owner might be 'working outside' to justify the *expectation*.

Both quantitative and qualitative analyses indicate that prompting methods can certainly complement, but may not fully replace, fine-tuned generation approaches to NORMBANK expansion. A mixed approach may be most desirable due to coverage and correctness, while generation errors may be fixed using self-correction via classification (above) or further prompting (Fung et al., 2022).

Finally, the middle pane of Table 3 shows the proportion of generated constraints that fall into our taxonomy (Tax. Constr.) and the proportion contained in NORMBANK (Pre-pop. Constr.). The former shows that our taxonomy broadly captures the relevant axes (80-90% of our best models' generations are taxonomic). The latter shows that between one third and one half of generations 'link' prior constraints to new situations; the rest of generated constraints are brand new.

5.2 Transfer Learning for Downstream Tasks

Finally, we conduct transfer learning experiments to demonstrate the utility of the data for downstream applications, further indicating the scope and power of NORMBANK as a general-purpose resource for social reasoning. Concretely, we follow the sequential training paradigm (Pratt et al., 1991), which has proven better than multitask training and fine-tuning on a broad range of commonsense tasks (Lourie et al., 2021a). Specifically, we initialize a RoBERTa model with weights from our best-performing norm classifier from Section 5.1 and fine-tune on the target set for 7 epochs.

We evaluate on two specifically *moral* reasoning tasks, Anecdotes and Dilemmas, both from the SCRUPLES benchmark (Lourie et al., 2021b). We also consider two multiple-choice commonsense QA datasets. Social IQa (Sap et al., 2019b) is designed to test social intelligence (e.g., inferring motivations, emotional reactions), while CosmosQA (Huang et al., 2019) tests cause and effect and counterfactual reasoning in everyday situations.

All results in Table 4 are averaged over five sepa-

Eval Task	Base Model	w/ Transfer Learning from		
		CQA	SIQA	NORMBANK
ANECDOTES	68.3	68.3‡	68.0	**68.7**⋆‡
DILEMMAS	64.3	67.4	70.9⋆	**71.1**⋆
SOCIALIQA	59.9	64.1	59.9	**64.2**
COSMOSQA	59.8	59.8	**63.5**⋆†	61.2

Table 4: **Transfer Learning Accuracies** demonstrate the utility of NORMBANK. By sequential finetuning on NORMBANK, we improve performance over baseline on all tasks, and transfer performance from NORMBANK exceeds transfer performance from CosmosQA and from SocialIQa in three cases. Best performance is **bolded**. Star⋆ results indicate significant improvements over the Base Model, while † marks significance over CQA, and ‡ marks significance over SIQA.

rate train-test runs, and significance is given by the paired bootstrap test. NORMBANK's utility is seen by comparing the accuracy of models with transfer learning from NORMBANK against those with task-only fine-tuning (Base Model). Results show that NORMBANK improves situational moral classification (Anecdotes; +0.4%) and forced choice binary moral judgments (Dilemmas +6.8%) with significance. Also consider NORMBANK utility as compared to transfer learning from either CosmosQA (CQA) or SocialIQa (SIQA). The only task on which transfer from NORMBANK does not achieve the best performance is on CosmosQA evaluation. Here, we find that transfer from the more structurally related Social IQa task is preferred. We conclude that NORMBANK is a useful resource for a range of downstream applications in moral, social, and emotional reasoning in context.

6 Conclusion

Social norms are the foundation of culture and society (McDonald and Crandall, 2015; Hogg and Reid, 2006), and an understanding of these norms is crucial for assistive and collaborative AI. In this work, we introduced SCENE a new scheme for hierarchically organizing the seemingly unbounded space of situational contexts that determine social norms. With this framework, we built NORMBANK, the first social knowledge bank to leverage such contextual information for contrast sets of richly conditioned defeasible social norms. We found that NORMBANK supports its own automatic expansion via classification, generation, and prompting methods. Finally, we demonstrated the utility of NORMBANK for situational social reasoning tasks.

7 Limitations

At its core, NORMBANK is a collection of logical operations on unique constraints. Consequently, one practical limitation stems from the issue that some situations cannot be reasonably expressed as a set of constraints. While theoretically all logic can be decomposed into AND and OR operations, the logic may be too challenging for an individual to formulate, or the set of constraints themselves might be too large and unwieldy. The latter is problematic, because language models have a finite input token capacity, and for the set of constraints to be digestible, they must fit within that capacity. Relatedly, if the logic to encode constraints become more sophisticated, ensuring that logic is not unnecessarily duplicated will pose a greater challenge. Additionally, certain properties of NORMBANK like the role and behaviors may be challenging to succinctly describe. Further work will be needed to ascertain how these can be incorporated or to more clearly define situations that are out of scope.

Due to limitations on time and computational resources, we have not exhaustively evaluated all downstream applications of NORMBANK, and in future work, we will test additional transfer tasks beyond the moral and social classification tasks considered in this work. Since NORMBANK is the first to encode non-monotonic situational norms, there was no other available *benchmark* that is directly analogous to ours. Instead, our primary evidence for NormBank's utility is in Table 3, where human evaluators confirm that models trained on NORMBANK can reliably learn to make new inferences about non-monotonic situational norms.

Other follow up studies should consider training larger normative reasoning models, and/or engineering better prompts for expanding NORMBANK. Relatedly, we have no data to speculate about the long-term evolution of real-world norms relative to this resource, nor the rate of decay in the reliability of NORMBANK. Future work should also expand this resource with perspectives from cultures other than our available annotator pool. The pool was not representative of all cultures and people groups, as we discuss further in the Ethics section.

8 Ethics

Ethical Assumptions. First, to set proper boundaries on this resource and the tasks it can facilitate, we will outline the ethical assumptions of this work and address some potential misconceptions. We want to stress that NORMBANK represents a collection of situational norms that we do not treat as prescriptive, but rather descriptive. Unlike prior *moral / ethics* datasets (Ziems et al., 2022; Emelin et al., 2021; Lourie et al., 2021b; Forbes et al., 2020; Sap et al., 2020), we use the neutral language of *expected, okay,* and *unexpected* behaviors to focus on empirically observed patterns and avoid an over-emphasis on the ethical grey area of what *ought* to be done. Unlike tricky moral dilemmas, the situational social norms of NORMBANK have an answer that a majority can agree is descriptively observable as the expectation under the respective conditions and/or cultural context. Nevertheless, normative judgments can vary between individuals in different social groups and time periods (Haidt et al., 1993; Shweder, 1990; Bicchieri, 2005; Culley and Madhavan, 2013; Amaya et al., 2021). NORM-BANK can and should be expanded via automatic or manual methods that can incorporate these axes of variation. Our annotator pool was limited to English-speaking individuals living in the United States in the year 2022. Future expansion efforts could be crowdsourced from other cultures and geographic regions and in future decades.

We reiterate that the norms in NORMBANK should *not* be used for prescriptive advice or personal guidance in any way. Our work intends to unlock future work in the capacity to imbue language models with situational commonsense and enable them to jointly reason with the situational contexts. Language models which ignore situational contexts altogether may be just as hazardous, if not more.

Finally, there are likely biases towards certain roles and values in NORMBANK. We have taken steps to mitigate some forms, such as gender bias, by neutralizing constraints (e.g., [PERSON]'s role is 'cowboy or cowgirl' and [PERSON]'s role is 'ball boy or girl'). Our SCENE taxonomy, with the standardized structure of its role and attribute constraints, will allow practitioners to further analyze specific axes of prejudice and thus implement targeted mitigation strategies. Specific identity attributes like gender, ethnicity, and religion are represented in 24% of norms.[5] Stakeholders can invest a smaller but more concerted effort towards mitigating bias in these constraints. We encourage

---

[5]There are 40k norms (out of 169k total norms; 24%) which cover attributes in the following set: {'country', 'age bracket', 'education', 'gender', 'race or ethnicity', 'religion', 'sexuality', 'social class'}

stakeholders to give auditing control over a given norm to those who are affected by it. Previous norm-datasets encode norms in free-text annotations which lack a hierarchical taxonomy of contexts, but our taxonomy can be used to interpret, diagnose, and mitigate prejudice, and to return power to those affected by these prejudices.

Risks in deployment. Before starting any annotation, the resources and findings presented in this work were thoroughly reviewed and approved by an internal review board. Prior to being put into production, the method would also need to be re-evaluated when applied to a new domain to ensure reliable performance in order to prevent unintended consequences. To help mitigate risks in deployment from misunderstandings about the ethical assumptions above, we require users of this data to complete a Data Use Agreement. The user will check that they understand the ethical assumptions above: especially that NORMBANK is not to be taken for advice. Practitioners will also agree not to use NORMBANK for malicious purposes "including (but not limited to): mockery, discrimination, and hate speech."

Acknowledgements

References

2022. Permanent missions to the united nations, no. 310. *Permanent Missions to the United Nations*.

Henk Aarts and Ap Dijksterhuis. 2003. The silence of the library: environment, situational norm, and social behavior. *Journal of personality and social psychology*, 84(1):18.

Irwin Altman. 1975. The environment and social behavior: privacy, personal space, territory, and crowding.

Ashley Amaya, Ruben Bach, Florian Keusch, and Frauke Kreuter. 2021. New data sources in social science research: things to know before working with reddit data. *Social science computer review*, 39(5):943–960.

Janna Anderson and Lee Rainie. 2022. The metaverse in 2040. *Pew Research Center*.

Rodrigo Bavaresco, Diórgenes Silveira, Eduardo Reis, Jorge Barbosa, Rodrigo Righi, Cristiano Costa, Rodolfo Antunes, Marcio Gomes, Clauter Gatti, Mariangela Vanzin, et al. 2020. Conversational agents in business: A systematic literature review and future research directions. *Computer Science Review*, 36:100239.

Paul A Bell, T Green, Jeffrey D Fisher, and Andrew Baum. 2001. *Environmental psychology*. New Jersey.

Cristina Bicchieri. 2005. *The grammar of society: The nature and dynamics of social norms*. Cambridge University Press.

Yonatan Bisk, Rowan Zellers, Ronan LeBras, Jianfeng Gao, and Yejin Choi. 2020. PIQA: reasoning about physical commonsense in natural language. In *The Thirty-Fourth AAAI Conference on Artificial Intelligence, AAAI 2020, The Thirty-Second Innovative Applications of Artificial Intelligence Conference, IAAI 2020, The Tenth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2020, New York, NY, USA, February 7-12, 2020*, pages 7432–7439. AAAI Press.

Joseph A Blass and Ian D Horswill. 2015. Implementing injunctive social norms using defeasible reasoning. In *Eleventh Artificial Intelligence and Interactive Digital Entertainment Conference*.

Michael Boratko, Xiang Li, Tim O'Gorman, Rajarshi Das, Dan Le, and Andrew McCallum. 2020. ProtoQA: A question answering dataset for prototypical common-sense reasoning. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1122–1136, Online. Association for Computational Linguistics.

Antoine Bordes, Nicolas Usunier, Alberto García-Durán, Jason Weston, and Oksana Yakhnenko. 2013. Translating embeddings for modeling multi-relational data. In *Advances in Neural Information Processing Systems 26: 27th Annual Conference on Neural Information Processing Systems 2013. Proceedings of a meeting held December 5-8, 2013, Lake Tahoe, Nevada, United States*, pages 2787–2795.

Peter R Boyce, Neil H Eklund, Barbara J Hamilton, and Lisa D Bruno. 2000. Perceptions of safety at night in different lighting conditions. *International Journal of Lighting Research and Technology*, 32(2):79–91.

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems 33:*

Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual.

Fabio Maria Carlucci, Lorenzo Nardi, Luca Iocchi, and Daniele Nardi. 2015. Explicit representation of social norms for social robots. In *2015 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 4191–4196. IEEE.

Jesse Chandler, Gabriele Paolacci, Eyal Peer, Pam Mueller, and Kate A Ratliff. 2015. Using nonnaive participants can reduce effect sizes. *Psychological science*, 26(7):1131–1139.

Ting-Yun Chang, Yang Liu, Karthik Gopalakrishnan, Behnam Hedayatnia, Pei Zhou, and Dilek Hakkani-Tur. 2020. Incorporating commonsense knowledge graph in pretrained models for social commonsense tasks. In *Proceedings of Deep Learning Inside Out (DeeLIO): The First Workshop on Knowledge Extraction and Integration for Deep Learning Architectures*, pages 74–79, Online. Association for Computational Linguistics.

Robert B Cialdini, Carl A Kallgren, and Raymond R Reno. 1991. A focus theory of normative conduct: A theoretical refinement and reevaluation of the role of norms in human behavior. In *Advances in experimental social psychology*, volume 24, pages 201–234. Elsevier.

Robert B Cialdini, Raymond R Reno, and Carl A Kallgren. 1990. A focus theory of normative conduct: Recycling the concept of norms to reduce littering in public places. *Journal of personality and social psychology*, 58(6):1015.

Mark Craven, Dan DiPasquo, Dayne Freitag, Andrew McCallum, Tom Mitchell, Kamal Nigam, and Seán Slattery. 2000. Learning to construct knowledge bases from the world wide web. *Artificial intelligence*, 118(1-2):69–113.

Kimberly E Culley and Poornima Madhavan. 2013. A note of caution regarding anthropomorphism in hci agents. *Computers in Human Behavior*, 29(3):577–579.

Michael R Cunningham. 1979. Weather, mood, and helping behavior: Quasi experiments with the sunshine samaritan. *Journal of personality and social psychology*, 37(11):1947.

Ken Dempsey and David De Vaus. 2004. Who cohabits in 2001? the significance of age, gender, religion and ethnicity. *Journal of Sociology*, 40(2):157–178.

Eric Deng, Bilge Mutlu, Maja J Mataric, et al. 2019. Embodiment in socially interactive robots. *Foundations and Trends® in Robotics*, 7(4):251–356.

Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.

E Durkheim. 1915. The elementary forms of the religious life: A study in religious sociology.

Hady Elsahar, Pavlos Vougiouklis, Arslen Remaci, Christophe Gravier, Jonathon Hare, Frederique Laforest, and Elena Simperl. 2018. T-REx: A large scale alignment of natural language with knowledge base triples. In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan. European Language Resources Association (ELRA).

Denis Emelin, Ronan Le Bras, Jena D Hwang, Maxwell Forbes, and Yejin Choi. 2021. Moral stories: Situated reasoning about norms, intents, actions, and their consequences. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 698–718.

Terrence Fong, Illah Nourbakhsh, and Kerstin Dautenhahn. 2003. A survey of socially interactive robots. *Robotics and autonomous systems*, 42(3-4):143–166.

Maxwell Forbes, Jena D. Hwang, Vered Shwartz, Maarten Sap, and Yejin Choi. 2020. Social chemistry 101: Learning to reason about social and moral norms. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 653–670, Online. Association for Computational Linguistics.

Yi R Fung, Tuhin Chakraborty, Hao Guo, Owen Rambow, Smaranda Muresan, and Heng Ji. 2022. Normsage: Multi-lingual multi-cultural norm discovery from conversations on-the-fly. *ArXiv preprint*, abs/2210.08604.

Matt Gardner, Yoav Artzi, Victoria Basmov, Jonathan Berant, Ben Bogin, Sihao Chen, Pradeep Dasigi, Dheeru Dua, Yanai Elazar, Ananth Gottumukkala, Nitish Gupta, Hannaneh Hajishirzi, Gabriel Ilharco, Daniel Khashabi, Kevin Lin, Jiangming Liu, Nelson F. Liu, Phoebe Mulcaire, Qiang Ning, Sameer Singh, Noah A. Smith, Sanjay Subramanian, Reut Tsarfaty, Eric Wallace, Ally Zhang, and Ben Zhou. 2020. Evaluating models' local decision boundaries via contrast sets. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 1307–1323, Online. Association for Computational Linguistics.

Ervin Goffman. 1959. *The presentation of self in everyday life*. Doubleday Anchor Books.

Joshua Grossman, Zhiyuan Lin, Hao Sheng, Johnny Tian-Zheng Wei, Joseph J Williams, and Sharad Goel. 2019. Mathbot: Transforming online resources for learning math into conversational interactions. *AAAI 2019 Story-Enabled Intelligence*.

Jonathan Haidt, Silvia Helena Koller, and Maria G Dias. 1993. Affect, culture, and morality, or is it wrong to eat your dog? *Journal of personality and social psychology*, 65(4):613.

Vicki S Helgeson. 2016. *Psychology of gender*. Routledge.

Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch Critch, Jerry Li Li, Dawn Song, and Jacob Steinhardt. 2021. Aligning ai with shared human values. In *International Conference on Learning Representations*.

Michael A Hogg and Scott A Reid. 2006. Social identity, self-categorization, and the communication of group norms. *Communication theory*, 16(1):7–30.

Dirk Hovy and Diyi Yang. 2021. The importance of modeling social factors of language: Theory and practice. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 588–602, Online. Association for Computational Linguistics.

Lifu Huang, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. 2019. Cosmos QA: Machine reading comprehension with contextual commonsense reasoning. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2391–2401, Hong Kong, China. Association for Computational Linguistics.

Gregory A Janicik and Caroline A Bartel. 2003. Talking about time: Effects of temporal planning and time awareness norms on group coordination and performance. *Group Dynamics: Theory, Research, and Practice*, 7(2):122.

Haozhe Ji, Pei Ke, Shaohan Huang, Furu Wei, and Minlie Huang. 2020a. Generating commonsense explanation by extracting bridge concepts from reasoning paths. In *Proceedings of the 1st Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 10th International Joint Conference on Natural Language Processing*, pages 248–257, Suzhou, China. Association for Computational Linguistics.

Haozhe Ji, Pei Ke, Shaohan Huang, Furu Wei, Xiaoyan Zhu, and Minlie Huang. 2020b. Language generation with multi-hop reasoning on commonsense knowledge graph. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 725–736, Online. Association for Computational Linguistics.

Liwei Jiang, Jena D Hwang, Chandra Bhagavatula, Ronan Le Bras, Maxwell Forbes, Jon Borchardt, Jenny Liang, Oren Etzioni, Maarten Sap, and Yejin Choi. 2021. Delphi: Towards machine ethics and norms. *ArXiv preprint*, abs/2110.07574.

Divyansh Kaushik, Eduard H. Hovy, and Zachary Chase Lipton. 2020. Learning the difference that makes A difference with counterfactually-augmented data. In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net.

Joanna Klimczyk. 2021. Compositional semantics and normative 'ought'. *Axiomathes*, 31(3):381–399.

Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. 2020. ALBERT: A lite BERT for self-supervised learning of language representations. In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net.

Douglas B Lenat. 1995. Cyc: A large-scale investment in knowledge infrastructure. *Communications of the ACM*, 38(11):33–38.

Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer. 2020. BART: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7871–7880, Online. Association for Computational Linguistics.

Yankai Lin, Zhiyuan Liu, Huanbo Luan, Maosong Sun, Siwei Rao, and Song Liu. 2015a. Modeling relation paths for representation learning of knowledge bases. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 705–714, Lisbon, Portugal. Association for Computational Linguistics.

Yankai Lin, Zhiyuan Liu, Maosong Sun, Yang Liu, and Xuan Zhu. 2015b. Learning entity and relation embeddings for knowledge graph completion. In *Proceedings of the Twenty-Ninth AAAI Conference on Artificial Intelligence, January 25-30, 2015, Austin, Texas, USA*, pages 2181–2187. AAAI Press.

Leib Litman, Jonathan Robinson, and Cheskie Rosenzweig. 2015. The relationship between motivation, monetary compensation, and data quality among us- and india-based workers on mechanical turk. *Behavior research methods*, 47(2):519–528.

Hugo Liu and Push Singh. 2004. Conceptnet—a practical commonsense reasoning tool-kit. *BT technology journal*, 22(4):211–226.

Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Roberta: A robustly optimized bert pretraining approach. *ArXiv preprint*, abs/1907.11692.

Ilya Loshchilov and Frank Hutter. 2019. Decoupled weight decay regularization. In *7th International Conference on Learning Representations, ICLR 2019,*

New Orleans, LA, USA, May 6-9, 2019. OpenReview.net.

Nicholas Lourie, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. 2021a. Unicorn on rainbow: A universal commonsense reasoning model on a new multitask benchmark. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 13480–13488.

Nicholas Lourie, Ronan Le Bras, and Yejin Choi. 2021b. Scruples: A corpus of community ethical judgments on 32, 000 real-life anecdotes. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 13470–13479.

Aman Madaan, Dheeraj Rajagopal, Niket Tandon, Yiming Yang, and Eduard Hovy. 2021. Could you give me a hint ? generating inference graphs for defeasible reasoning. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 5138–5147, Online. Association for Computational Linguistics.

Kenneth E Mathews and Lance Kirkpatrick Canon. 1975. Environmental noise level as a determinant of helping behavior. *Journal of Personality and Social Psychology*, 32(4):571.

Rachel I McDonald and Christian S Crandall. 2015. Social norms and social influence. *Current Opinion in Behavioral Sciences*, 3:147–151.

Erin Meyer. 2014. *The culture map: Breaking through the invisible boundaries of global business*. Public Affairs.

Tom M. Mitchell, William W. Cohen, Estevam R. Hruschka Jr., Partha Pratim Talukdar, Justin Betteridge, Andrew Carlson, Bhavana Dalvi Mishra, Matthew Gardner, Bryan Kisiel, Jayant Krishnamurthy, Ni Lao, Kathryn Mazaitis, Thahir Mohamed, Ndapandula Nakashole, Emmanouil A. Platanios, Alan Ritter, Mehdi Samadi, Burr Settles, Richard C. Wang, Derry Wijaya, Abhinav Gupta, Xinlei Chen, Abulhair Saparov, Malcolm Greaves, and Joel Welling. 2015. Never-ending learning. In *Proceedings of the Twenty-Ninth AAAI Conference on Artificial Intelligence, January 25-30, 2015, Austin, Texas, USA*, pages 2302–2310. AAAI Press.

Arindam Mitra, Pratyay Banerjee, Kuntal Kumar Pal, Swaroop Mishra, and Chitta Baral. 2019. How additional knowledge can improve natural language commonsense question answering? *ArXiv preprint*, abs/1909.08855.

György Molnár and Zoltán Szüts. 2018. The role of chatbots in formal education. In *2018 IEEE 16th International Symposium on Intelligent Systems and Informatics (SISY)*, pages 000197–000202. IEEE.

Fabio Petroni, Tim Rocktäschel, Sebastian Riedel, Patrick Lewis, Anton Bakhtin, Yuxiang Wu, and Alexander Miller. 2019. Language models as knowledge bases? In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2463–2473, Hong Kong, China. Association for Computational Linguistics.

John L Pollock. 1987. Defeasible reasoning. *Cognitive science*, 11(4):481–518.

Lorien Y Pratt, Jack Mostow, Candace A Kamm, Ace A Kamm, et al. 1991. Direct transfer of learned information among neural networks. In *Aaai*, volume 91, pages 584–589.

Valentina Pyatkin, Jena D Hwang, Vivek Srikumar, Ximing Lu, Liwei Jiang, Yejin Choi, and Chandra Bhagavatula. 2022. Reinforced clarification question generation with defeasibility rewards for disambiguating social and moral situations. *ArXiv preprint*, abs/2212.10409.

Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9.

Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. 2020. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67.

Raymond Reiter. 1981. On closed world data bases. In *Readings in artificial intelligence*, pages 119–140. Elsevier.

Rachel Rudinger, Vered Shwartz, Jena D. Hwang, Chandra Bhagavatula, Maxwell Forbes, Ronan Le Bras, Noah A. Smith, and Yejin Choi. 2020. Thinking like a skeptic: Defeasible inference in natural language. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 4661–4675, Online. Association for Computational Linguistics.

James A Russell and Lawrence M Ward. 1982. Environmental psychology. *Annual review of psychology*, 33(1):651–689.

Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. 2020. Winogrande: An adversarial winograd schema challenge at scale. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 8732–8740.

Maarten Sap, Ronan Le Bras, Emily Allaway, Chandra Bhagavatula, Nicholas Lourie, Hannah Rashkin, Brendan Roof, Noah A. Smith, and Yejin Choi. 2019a. ATOMIC: an atlas of machine commonsense for if-then reasoning. In *The Thirty-Third AAAI Conference on Artificial Intelligence, AAAI 2019, The Thirty-First Innovative Applications of Artificial Intelligence Conference, IAAI 2019, The Ninth AAAI Symposium on Educational Advances in Artificial*

Intelligence, EAAI 2019, Honolulu, Hawaii, USA, January 27 - February 1, 2019, pages 3027–3035. AAAI Press.

Maarten Sap, Saadia Gabriel, Lianhui Qin, Dan Jurafsky, Noah A. Smith, and Yejin Choi. 2020. Social bias frames: Reasoning about social and power implications of language. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5477–5490, Online. Association for Computational Linguistics.

Maarten Sap, Hannah Rashkin, Derek Chen, Ronan Le Bras, and Yejin Choi. 2019b. Social IQa: Commonsense reasoning about social interactions. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 4463–4473, Hong Kong, China. Association for Computational Linguistics.

R. C. Schank and R. P. Abelson. 1977. *Scripts, plans, goals, and understanding: An enquiry into human knowledge structures*. Erlbaum.

Alister Scott, Alana Gilbert, and Ayele Gelan. 2007. *The urban-rural divide: Myth or reality?* Macaulay Institute Aberdeen.

Kim Bartel Sheehan. 2018. Crowdsourcing research: data collection with amazon's mechanical turk. *Communication Monographs*, 85(1):140–156.

Muzafer Sherif and Carolyn W Sherif. 1953. Groups in harmony and tension; an integration of studies of intergroup relations.

Baoxu Shi and Tim Weninger. 2018. Open-world knowledge graph completion. In *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence, (AAAI-18), the 30th innovative Applications of Artificial Intelligence (IAAI-18), and the 8th AAAI Symposium on Educational Advances in Artificial Intelligence (EAAI-18), New Orleans, Louisiana, USA, February 2-7, 2018*, pages 1957–1964. AAAI Press.

Richard A Shweder. 1990. In defense of moral realism: Reply to gabennesch. *Child Development*, 61(6):2060–2067.

Richard Socher, Danqi Chen, Christopher D. Manning, and Andrew Y. Ng. 2013. Reasoning with neural tensor networks for knowledge base completion. In *Advances in Neural Information Processing Systems 26: 27th Annual Conference on Neural Information Processing Systems 2013. Proceedings of a meeting held December 5-8, 2013, Lake Tahoe, Nevada, United States*, pages 926–934.

Robyn Speer, Joshua Chin, and Catherine Havasi. 2017. Conceptnet 5.5: An open multilingual graph of general knowledge. In *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence, February 4-9, 2017, San Francisco, California, USA*, pages 4444–4451. AAAI Press.

Cass R Sunstein. 1996. Social norms and social roles. *Colum. L. Rev.*, 96:903.

Zeerak Talat, Hagen Blix, Josef Valvoda, Maya Indira Ganesh, Ryan Cotterell, and Adina Williams. 2022. On the machine learning of ethical judgments from natural language. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 769–779, Seattle, United States. Association for Computational Linguistics.

Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. 2019. CommonsenseQA: A question answering challenge targeting commonsense knowledge. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4149–4158, Minneapolis, Minnesota. Association for Computational Linguistics.

Edward H Thompson Jr and Joseph H Pleck. 1986. The structure of male role norms. *American Behavioral Scientist*, 29(5):531–543.

Trieu H Trinh and Quoc V Le. 2018. A simple method for commonsense reasoning. *ArXiv preprint*, abs/1806.02847.

Aditya Nrusimha Vaidyam, Hannah Wisniewski, John David Halamka, Matcheri S Kashavan, and John Blake Torous. 2019. Chatbots and conversational agents in mental health: a review of the psychiatric landscape. *The Canadian Journal of Psychiatry*, 64(7):456–464.

Leon van Rijswijk and Antal Haans. 2018. Illuminating for safety: Investigating the role of lighting appraisals on the perception of safety in the urban environment. *Environment and behavior*, 50(8):889–912.

Paul Vicol, Makarand Tapaswi, Lluís Castrejón, and Sanja Fidler. 2018. Moviegraphs: Towards understanding human-centric situations from videos. In *2018 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2018, Salt Lake City, UT, USA, June 18-22, 2018*, pages 8581–8590. IEEE Computer Society.

Iris Vilnai-Yavetz and Shaked Gilboa. 2010. The effect of servicescape cleanliness on customer reactions. *Services Marketing Quarterly*, 31(2):213–234.

Cunxiang Wang, Shuailong Liang, Yue Zhang, Xiaonan Li, and Tian Gao. 2019. Does it make sense? and why? a pilot study for sense making and explanation. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4020–4026, Florence, Italy. Association for Computational Linguistics.

Jason Weston, Antoine Bordes, Oksana Yakhnenko, and Nicolas Usunier. 2013. Connecting language and knowledge bases with embedding models for relation extraction. In *Proceedings of the 2013 Conference*

on Empirical Methods in Natural Language Processing, pages 1366–1371, Seattle, Washington, USA. Association for Computational Linguistics.

Sheng Zhang, Xiaodong Liu, Jingjing Liu, Jianfeng Gao, Kevin Duh, and Benjamin Van Durme. 2018. Record: Bridging the gap between human and machine commonsense reading comprehension.

Caleb Ziems, Jane A. Yu, Yi-Chia Wang, Alon Y. Halevy, and Diyi Yang. 2022. The moral integrity corpus: A benchmark for ethical dialogue systems. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing*, Dublin, Ireland. Association for Computational Linguistics.

A Models & Hyperparameters

Classification.   We use the base versions of BERT (Devlin et al.; 768-hidden, 12-heads, 110M parameters), RoBERTa (Liu et al.; 768-hidden, 12-heads, 125M parameters), and ALBERT-v2 (Lan et al.; 768-hidden, 12-heads, 11M parameters). For each model, we fine-tune using AdamW (Loshchilov and Hutter, 2019) for 7 epochs with a batch size of 16 and a learning rate of $1e-5$. These hyperparameters were chosen by hyperparameter search on the dev set over $\{1e-5, 2e-5, 3e-5, 5e-5\}$ and the number of epochs in $\{1..8\}$, with $\epsilon = 1e-8$ and the batch size set to 16.

Generation.   We trained BART-large (406M parameters), GPT-2 (768-hidden, 12-heads, 117M parameters), and T5-small (512-hidden, 8 heads, 60M parameters) for 1 epoch using a batch size of 8 and a learning rate of $3e-5$. We also prompted GPT-3 Davinci-002 and Davinci-003 in a few-shot manner via the OpenAI API, using the following prompt, which contains seven representative datapoints from the NORMBANK train set:

```
  It is UNEXPECTED to "talk about sex" at an "cafe" when:
+ PERSON's role is 'barista'

It is OKAY to "listen to music" in an "arcade" when all of
    ↪ the following are true:
+ PERSON's role is not 'cook'
+ PERSON's role is not 'maintenance worker'
+ PERSON's behavior is 'wear headphones'

It is EXPECTED to "hang film to dry" at the "darkroom" when
    ↪ all of the following are true:
+ PERSON's role is 'darkroom technician'
+ PERSON's behavior is 'performing routine tasks'

It is OKAY to "smoke a cigar" in a "limousine" when all of
    ↪ the following are true:
+ OTHER's age bracket is not 'gradeschooler or younger'
+ OTHER's attitude towards smoking is not 'bothered by smoke'
    ↪
+ PERSON's age bracket is 'adult'
+ limousine rules is not 'forbids smoking'

It is EXPECTED to "maintain the store appearance" at the "
    ↪ clothing store" when all of the following are true:
+ PERSON's role is 'store owner'

It is OKAY to "talk to yourself" on an "airplane" when all
    ↪ of the following are true:
+ PERSON's role is 'passenger'
+ PERSON's behavior is 'pray'

It is UNEXPECTED to "grow tomatoes" in a "garden" when all
    ↪ of the following are true:
+ PERSON's behavior is not 'use a greenhouse'
+ temperature is 'freezing'

It is {norm} to "{behavior}" {prep} {det} "{setting}" when
    ↪ all of the following are true:
+
```

B Additional Details on Constructing SCENE

B.1 The Environment 🏞️

Country of Operation   is seeded with the 195 countries from the UN (2022) list of member or non-member observer states.

Operational Factors   is a broad category of constraints from environmental psychology (Bell et al., 2001) involving one's comfort and ease of operation in an environment. Such factors influence descriptive norms. Operational behaviors can be influenced by the degree of sensory stimulation, as well as by privacy and proxemics, or the local density and organization of persons and objects (Russell and Ward, 1982). These inform the following subcategories: **attendance** {*empty, there are people around, crowded*} (Altman, 1975); **cleanliness** {*dirty, clean*} (Vilnai-Yavetz and Gilboa, 2010; Cialdini et al., 1990); **noise** {*quiet, moderate, loud*} (Mathews and Canon, 1975); **population density** {*urban, suburban, rural*} (Scott et al., 2007); and **privacy** {*private, public*} (Altman, 1975).

Physical Conditions   in the environment can influence behavior mechanically as well as psychologically. Specifically, the **lighting** {*bright, moderate, dim, dark*}, **weather** {*blizzard, clear, cloudy, ...*}, and **temperature** {*freezing, cold, temperate, hot*} can directly impact visibility, coordination, and the perception of safety (Boyce et al., 2000; van Rijswijk and Haans, 2018), as well as comfort, confidence, and altruism (Cunningham, 1979).

Restrictions   formally limit attendance, participation, and behavior. The environment can be one of **exclusion** {*adults only, men only, women only*}, **formality** {*formal, informal*} or **religiosity** {*sacred, secular*}. These categories are not part of our original theoretical taxonomy but were introduced through annotator feedback (See Section 4).

Special Observances   include cultural observances like **holidays** {*Advent, Holi, Lunar New Year*} as well other **special events** {*bat mitzvah, housewarming, quinceañera*}, which evoke distinct rituals, customs and norms (Durkheim, 1915).

Time   constraints include **day of the week**, **season**, **time of day**, and **time period**. Like special observances, much of human activity adheres to a set of temporal constraints and cues (Janicik and Bartel, 2003).

B.2 Behaviors 🏸

Since the average string length of behaviors was greater and thus more prone to error, we applied a suite of programmatic cleaning and filtering techniques, followed by a manual filtering round that reduced the average to 112.6 clean and non-

redundant behaviors per setting. The filtering techniques are as follows:

1. Remove any conditional form "if you were..." as well as any mention of the role or the setting in the behavior itself

2. Remove elaborations on behaviors "<behavior> because ... <elaboration>"

3. Normalize the logical form by removing words that negate behaviors (do not, never, not, forget to, refuse to, fail to, anything, something, anyone, someone, in any way, any)

4. Remove biased terms like *properly, should, would, try to, be able to, and see someone*.

5. Remove any bullet points or stray characters

6. Require that the behavior has an active (not passive) verb in it and that the verb does not have an explicit subject and that there is no dependent clause (as indicated by the marker `[mark]` dependency).

C Annotation Task Details

C.1 Qualification Task

To qualify for the HIT, workers were required to pass the following qualifying test, answering at least 5 out of 6 questions correctly.

1. **True or False:** you are allowed to add your own Constraints by typing them directly into the box. [*Answer*: **True**]

2. **True or False:** "carry a gun" can be an "Expected" behavior on an Airplane. [*Answer*: **True**]

3. **True or False:** "read a book" is an "Expected" behavior for a passenger on an Airplane. [*Answer*: **False**]

4. **True or False:** it is possible for a BEHAVIOR to be both "Expected" and "Okay" under the same Constraints. [*Answer*: **False**]

5. Let's say you are adding some Constraints for when "eating shrimp" is "Unexpected" in the SETTING: restaurant. You know that shellfish are disallowed in both Hinduism and Judaism, as well as by the vegan and vegetarian diets. You are thinking of adding these in the following Constraint table. Is this correct? [*Answer*: **Incorrect**]

(PERSON's religion is Judaism) AND (PERSON's religion is Hinduism) AND (PERSON's diet is vegan) AND (PERSON's diet is vegetarian)

6. Let's say you are adding some Constraints for when "eating shrimp" is "Okay" in the SETTING: restaurant. You know that shellfish are disallowed in both Hinduism and Judaism, as well as by the vegan and vegetarian diets. You are thinking of adding these in the following Constraint table. Is this correct? [*Answer*: **Correct**]

(PERSON's religion is NOT Judaism) AND (PERSON's religion is NOT Hinduism) AND (PERSON's diet is NOT vegan) AND (PERSON's diet is NOT vegetarian)

C.2 HIT Interface

For each HIT, the annotator is presented with a setting $s \in \mathcal{S}$ and a behavior $b \in \mathcal{B}$ that we generated for the given s. The annotator helps us describe when this behavior would be *expected*; then describes when it is merely *okay*; and finally *unexpected*. Annotators describe each norm with the conjunction and disjunction of SCENE constraints. The annotator appends each constraint to its conjunction as a 4-tuple consisting of a (1) *category*, (2) *name*, (3) *relation*, and (4) *value*. These are shown with examples in the HIT Instructions (4) and HIT Interface (Figure 5) screenshots. The *category* helps annotators search for constraints and organize their thoughts. The *category* is a high-level designation of where the constraint is organized: according to the *environment*, *role*, *attribute*, or *behavior*. The *name, relation* and *value* constitute a standard semantic triple. The *name* designates the subject of the constraint, and it a specification of the *category*, like the "temperature of the environment." The *relation* is a logical type that includes equality and inequality. The *value* designates the predicate of the constraint (e.g., "freezing").

Annotators can build constraint 4-tuples from drop-down menus that enumerate our hierarchical taxonomy (Section 3). Annotators can also freely edit the above fields and contribute novel constraints. Finally, annotators compose constraints into disjunctive normal form (DNF), the OR of ANDs, to describe when behaviors are *expected, okay,* or *unexpected* in a given setting.

Instructions   (Expand/Collapse)

In this HIT, you will be given the name of a social SETTING like an *airplane*. Then you can help think of examples of different Constraints where the BEHAVIOR might be expected, okay, or unexpected.

By "expected, okay, or unexpected" we use the following definitions.

1. **Expected:** Behaviors that a certain kind of person would be predicted or required to do in this setting. There is often some kind of *obligation* here.
2. **Okay:** Behaviors that a certain kind of person would reasonably do in this setting, but there is *no obligation*.
3. **Unexpected:** Behaviors that may be socially or culturally taboo, forbidden, discouraged, or generally surprising and "not normal" for a certain kind of person to do in this setting.

There are examples of Constraints in the tables below. Constraints have a Category, Name, Relation, and Value. Play around with these options to understand them better.

- You should think of the Constraints in a single table as having the word AND between them. That is, all the Constraints work together to define the conditions where the BEHAVIOR might be expected, okay, or unexpected.
- The PERSON PERFORMING BEHAVIOR is main person you are adding Constraints for, but you can also add Constraints for some OTHER PERSON. For example, if the BEHAVIOR is "serving drinks" and the SETTING is "airplane," the PERSON PERFORMING BEHAVIOR is often the flight attendant, and the OTHER PERSON is the passenger who is receiving the drink.
- If the default list doesn't contain the Constraint that you need, feel free to create a new Constraint by directly typing it into the box.

Please review these examples. You can hover over any highlighted text in the instructions for important additional insights.

Examples of Expected Behaviors
(under the following Constraints)

SETTING: *airplane*

- **BEHAVIOR:** *communicate with air traffic control* **is Expected when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's role	airplane role	is	copilot	By definition of her job, the copilot is expected to communicate with air traffic control.

- **BEHAVIOR:** *distribute snacks* **is Expected when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's role	airplane role	is	flight attendant	As part of their job, flight attendants are expected to distribute snacks.
2.	some OTHER PERSON's role	airplane role	is	passenger	Snacks should be distributed only to passengers.
3.	some OTHER PERSON's attribute	state	is NOT	sleeping	Sleeping passengers should not be woken up and offered snacks.

- **BEHAVIOR:** *bring a weapon* **is Expected when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's role	airplane role	is	air marshal	An air marshal is a federal agent disguised to look like regular passenger. Each air marshal is authorized and expected to carry a gun for emergencies.

Examples of Okay Behaviors
(under the following Constraints)

SETTING: *airplane*

- **BEHAVIOR:** *eat a ham sandwich* **is okay when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's role	airplane role	is	passenger	While other roles like pilots and flight assistants might not, passengers often eat complimentary snacks on a flight.
2.	the PERSON PERFORMING BEHAVIOR's attribute	age-bracket	is NOT	infant	Infants could choke on a sandwich.
3.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Judaism	Pork is not kosher in Judaism.
4.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Islam	Pork is haram (forbidden) in Islam.
5.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Seventh Day Adventist	Adventists recognize pork as unclean.

- **BEHAVIOR:** *chew on the tray table* **is Okay when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's attribute	age-bracket	is	infant	While it certainly is not required, infants are known to chew on things like tray tables. This would not be considered taboo for an infant to do.

SETTING: *boat*

- **BEHAVIOR:** *drink beer* **is Okay when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's attribute	country	is	the United States of America	We set the country first because this determines the drinking age.
2.	the PERSON PERFORMING BEHAVIOR's attribute	age	is greater than or equal to	21	This is the drinking age in the US.
3.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Islam	Muslims do not drink alcohol.
4.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Jainism	Jains do not drink alcohol.
5.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Bahá'í	Bahá'ís do not drink alcohol.
6.	the PERSON PERFORMING BEHAVIOR's role	boat role	is NOT	captain	If this person is the captain or boat operator, they must remain sober or else they could face legal trouble for Boating Under the Influence (BUI).

Examples of Okay Behaviors
(under the following Constraints)

SETTING: *airplane*

- **BEHAVIOR:** *eat a ham sandwich* **is okay when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's role	airplane role	is	passenger	While other roles like pilots and flight assistants might not, passengers often eat complimentary snacks on a flight.
2.	the PERSON PERFORMING BEHAVIOR's attribute	age-bracket	is NOT	infant	Infants could choke on a sandwich.
3.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Judaism	Pork is not kosher in Judaism.
4.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Islam	Pork is haram (forbidden) in Islam.
5.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Seventh Day Adventist	Adventists recognize pork as unclean.

- **BEHAVIOR:** *chew on the tray table* **is Okay when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's attribute	age-bracket	is	infant	While it certainly is not required, infants are known to chew on things like tray tables. This would not be considered taboo for an infant to do.

SETTING: *boat*

- **BEHAVIOR:** *drink beer* **is Okay when...**

Constraints

	Category	Name	Relation	Value	Explanation
1.	the PERSON PERFORMING BEHAVIOR's attribute	country	is	the United States of America	We set the country first because this determines the drinking age.
2.	the PERSON PERFORMING BEHAVIOR's attribute	age	is greater than or equal to	21	This is the drinking age in the US.
3.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Islam	Muslims do not drink alcohol.
4.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Jainism	Jains do not drink alcohol.
5.	the PERSON PERFORMING BEHAVIOR's attribute	religion	is NOT	Bahá'í	Bahá'ís do not drink alcohol.
6.	the PERSON PERFORMING BEHAVIOR's role	boat role	is NOT	captain	If this person is the captain or boat operator, they must remain sober or else they could face legal trouble for Boating Under the Influence (BUI).

Figure 4: HIT Instructions.

Figure 5: HIT Interface